

# PushNet: Efficient and Adaptive Neural Message Passing

Julian Busch <sup>12</sup>, Jiaxing Pi <sup>1</sup> and Thomas Seidl <sup>2</sup>


<sup>1</sup> Siemens Corporate Technology, Princeton, NJ, USA  
jiaxing.pi@siemens.com

<sup>2</sup> Ludwig-Maximilians-Universität München, Munich, Germany  
{busch, seidl}@dbs.ifi.lmu.de






# Introduction

- ▶ Semi-supervised node classification
  - ▶ Given
    - ▶ Graph  $G = (V, E)$
    - ▶ Feature matrix  $X \in \mathbb{R}^{n \times d}$
    - ▶ Incomplete label matrix  $Y \in \mathbb{R}^{n \times c}$
  - ▶ Goal
    - ▶ Predict labels of unlabeled nodes
- ▶ Applications
  - ▶ Document classification in citation networks
  - ▶ User recommendations in social networks
  - ▶ Function prediction in protein interaction networks
  - ▶ And many more!



# Neural Message Passing [1–15]

- ▶ Nodes repeatedly pull features from their neighbors:
  1. *Send messages* to all neighbors
  2. *Aggregate* incoming messages
  3. *Update* own feature vector
- ▶  $k$  rounds of message passing gathers features from  $k$ -hop neighborhood
- ▶ A node's label is predicted from its updated feature vector




# Motivation

- ▶ These methods are very successful
- ▶ However, they share several severe issues:
  - ▶ Not adaptive: In each iteration, all nodes send messages to all neighbors
    - ▶ Over-smoothing effect for multi-layer message passing [16, 17]
    - ▶ Inefficient: Only a few long-range dependencies are actually relevant
  - ▶ Restriction to  $k$ -hop neighborhoods
    - ▶ Include irrelevant, miss relevant neighbors
    - ▶ Need to specify number of message passing rounds  $k$

*Idea: Push information on demand instead of just pulling it from all neighbors!*

# Push-based Asynchronous Message Passing

- ▶ Main idea
  - ▶ Targeted propagation by updating nodes sequentially
- ▶ General procedure
  - ▶ Each node aggregates incoming messages until it is chosen to be updated
  - ▶ The chosen node updates its own state and pushes messages to all its neighbors
  - ▶ A node is only considered for updating if it aggregated enough new information



*Challenge: Can we do this and still be efficient?*


# PushNet

- *Yes!* PushNet can be formulated as a single synchronous message passing step:

$$H^{(0)} = f(X; \theta_f) \quad \in \mathbb{R}^{n \times h_1}$$

$$H^{(1)} = PH^{(0)} \quad \in \mathbb{R}^{n \times h_1}$$

$$H^{(2)} = g(H^{(1)}; \theta_g) \quad \in \mathbb{R}^{n \times h_2}$$



- The weights  $p_{ij}$  are given by *Approximate Personalized PageRank (APPR)* [18–20]
- $g$  and  $f$  are MLPs and act as learnable feature transformations/predictor

*Combines the best of both worlds of synchronous/asynchronous message passing!*

# Extensions

- ▶ Multi-scale neighborhood aggregation
  - ▶ Locality hyper-parameter  $\alpha$  in *APPR* controls the effective neighborhood size
  - ▶ Idea: Propagate over multiple scales instead of just one and aggregate the results
  - ▶ Benefits:
    - ▶ *Simplified hyper-parameter search*: Just combine a set of potential values
    - ▶ We empirically observe *improved classification accuracy and robustness*
- ▶ Model variants
  - ▶ Learnable feature transformations could be applied before and after propagation
  - ▶ General model: *PushNet*
  - ▶ No transformations before propagation
    - ▶ *Fast, propagated features can be cached!*
    - ▶ *PushNet-PTP, PushNet-PP*
  - ▶ Propagation of predicted class labels
    - ▶ *PushNet-TPP*

# Experimental Setup

- ▶ Evaluation on 5 document classification benchmark datasets
- ▶ Comparison with 7 SOTA neural message passing algorithms
- ▶ Rigorous evaluation setup [21, 22]
  - ▶ Grid search to determine best hyper-parameters for each model
  - ▶ Results from 100 independent runs

	$ V $	$ E $	$d$	$c$	avgSP	maxSP
<b>CiteSeer</b>	2120	3679	3703	6	9.33	28
<b>Cora</b>	2485	5069	1433	7	6.31	19
<b>PubMed</b>	19717	44324	500	3	6.34	18
<b>Coauthor CS</b>	18333	81894	6805	15	5.43	24
<b>Coauthor Physics</b>	34493	247962	8415	5	5.16	17




# Node Classification Accuracy

	CiteSeer	Cora	PubMed	Coauthor CS	Coauthor Physics
<b>GCN</b> [4]	$72.82 \pm 1.48$	$81.07 \pm 1.43$	$78.29 \pm 1.48$	$91.64 \pm 0.62$	$93.42 \pm 0.63$
<b>GAT</b> [7]	$73.82 \pm 1.35$	$82.12 \pm 1.41$	$78.21 \pm 1.60$	$90.20 \pm 0.75$	$93.43 \pm 0.50$
<b>JK-GCN</b> [17]	$71.09 \pm 1.66$	$79.57 \pm 1.63$	$77.23 \pm 2.01$	$91.60 \pm 0.54$	$93.49 \pm 0.56$
<b>JK-GAT</b> [17]	$71.76 \pm 1.27$	$80.10 \pm 1.52$	$77.59 \pm 2.25$	$92.20 \pm 0.43$	<i>o.o.m.</i>
<b>SGC</b> [15]	$73.91 \pm 1.30$	$80.13 \pm 2.15$	$77.00 \pm 1.78$	$91.27 \pm 0.58$	<i>o.o.m.</i>
<b>GIN</b> [13]	$70.81 \pm 1.61$	$80.24 \pm 1.54$	$77.19 \pm 1.75$	$91.46 \pm 0.54$	$93.79 \pm 0.49$
<b>APPNP</b> [22]	$74.36 \pm 1.44$	$83.58 \pm 1.03$	$79.61 \pm 2.98$	$91.10 \pm 1.12$	$93.96 \pm 0.45$
<b>PushNet</b>	$75.08 \pm 0.99$	$84.12 \pm 1.08$	$79.80 \pm 1.39$	$92.40 \pm 0.52$	$94.01 \pm 0.53$
<b>PushNet-PTP</b>	<b><math>75.19 \pm 1.15</math></b>	$83.41 \pm 1.24$	<b><math>80.22 \pm 1.27</math></b>	$92.37 \pm 0.40$	$93.97 \pm 0.48$
<b>PushNet-PP</b>	$75.17 \pm 1.32$	$81.52 \pm 1.40$	$77.52 \pm 2.05$	$91.04 \pm 0.76$	$93.67 \pm 0.55$
<b>PushNet-TPP</b>	$75.01 \pm 1.11$	<b><math>84.23 \pm 1.26</math></b>	$80.10 \pm 1.33$	<b><math>92.54 \pm 0.34</math></b>	<b><math>94.09 \pm 0.47</math></b>

# Node Classification Accuracy


- ▶ *PushNet* already outperforms all competitors on all datasets
  - ▶ *PushNet-TPP* performs best overall
  - ▶ *PushNet-PTP* improves performance on *CiteSeer* and *PubMed*
  - ▶ *PushNet-PP* is remarkably competitive, even though it performs no learnable feature transformation at all!
- ▶ All improvements are statistically significant
- ▶ Improvements due to *individually adapted neighborhoods for each node*



Fraction of  $k$ -neighbors included in *APPR*-neighborhood for each node

- ▶ Solid: Average over all nodes
- ▶ Red: 10-neighborhood



# Runtime



- ▶ *SGC* [15] is fastest but less accurate than most competitors
- ▶ *PushNet-PP* is second fastest, followed by *PushNet-PTP* which offers a good tradeoff between runtime and accuracy
- ▶ *PushNet* and *PushNet-TTP* are slower but provide superior accuracy

# Ablation Study

- ▶ Best single scales:  $\alpha \in \{0.05, 0.1, 0.2\}$
- ▶ Best aggregation: *sum*
  - ▶ Efficient: Runtime determined by smallest single  $\alpha$  considered
  - ▶ Improved results over best single  $\alpha$
- ▶ Sparsity
  - ▶ Parameter  $\varepsilon$  in *APPR* to exclude small propagation weights
  - ▶ Accuracy remains very stable
  - ▶ Intuitive to set, just set large enough for available GPU memory!



# Conclusion

- ▶ Novel asynchronous approach to neural message passing
- ▶ Efficient aggregation over adaptive node neighborhoods
  - ▶ No restriction to  $k$ -neighborhoods
  - ▶ Effective handling of long-range dependencies
- ▶ Multi-scale feature aggregation
  - ▶ Improved performance
  - ▶ Simplified hyper-parameter selection
- ▶ Different variants of our base model
  - ▶ Allow for varying trade-offs between efficiency and accuracy
- ▶ Hyper-parameters relatively easy to set

Thank you! Questions?

# References I

- [1] David K Duvenaud, Dougal Maclaurin, Jorge Iparraguirre, Rafael Bombarell, Timothy Hirzel, Alán Aspuru-Guzik, and Ryan P Adams. Convolutional networks on graphs for learning molecular fingerprints. *NeurIPS*, 2015.
- [2] Yujia Li, Daniel Tarlow, Marc Brockschmidt, and Richard Zemel. Gated graph sequence neural networks. *ICLR*, 2016.
- [3] Steven Kearnes, Kevin McCloskey, Marc Berndl, Vijay Pande, and Patrick Riley. Molecular graph convolutions: moving beyond fingerprints. *Journal of computer-aided molecular design*, 2016.
- [4] Thomas N Kipf and Max Welling. Semi-supervised classification with graph convolutional networks. *ICLR*, 2017.
- [5] Will Hamilton, Zhitaoy Ying, and Jure Leskovec. Inductive representation learning on large graphs. *NeurIPS*, 2017.
- [6] Justin Gilmer, Samuel S Schoenholz, Patrick F Riley, Oriol Vinyals, and George E Dahl. Neural message passing for quantum chemistry. *ICML*, 2017.

## References II

- [7] Petar Veličković, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Lio, and Yoshua Bengio. Graph attention networks. *ICLR*, 2018.
- [8] Jie Chen, Tengfei Ma, and Cao Xiao. Fastgcn: fast learning with graph convolutional networks via importance sampling. *ICLR*, 2018.
- [9] John Boaz Lee, Ryan A Rossi, Sungchul Kim, Nesreen K Ahmed, and Eunyee Koh. Attention models in graphs: A survey. *TKDD*, 2019.
- [10] Christopher Morris, Martin Ritzert, Matthias Fey, William L Hamilton, Jan Eric Lenssen, Gaurav Rattan, and Martin Grohe. Weisfeiler and leman go neural: Higher-order graph neural networks. *AAAI*, 2019.
- [11] Michael Schlichtkrull, Thomas N Kipf, Peter Bloem, Rianne Van Den Berg, Ivan Titov, and Max Welling. Modeling relational data with graph convolutional networks. *ESWC*, 2018.



## References III

- [12] Kiran K Thekumparampil, Chong Wang, Sewoong Oh, and Li-Jia Li. Attention-based graph neural network for semi-supervised learning. *arXiv preprint arXiv:1803.03735*, 2018.
- [13] Keyulu Xu, Weihua Hu, Jure Leskovec, and Stefanie Jegelka. How powerful are graph neural networks? *ICLR*, 2019.
- [14] Matthias Fey and Jan Eric Lenssen. Fast graph representation learning with pytorch geometric. *arXiv preprint arXiv:1903.02428*, 2019.
- [15] Felix Wu, Tianyi Zhang, Amauri Holanda de Souza Jr, Christopher Fifty, Tao Yu, and Kilian Q Weinberger. Simplifying graph convolutional networks. *ICML*, 2019.
- [16] Qimai Li, Zhichao Han, and Xiao-Ming Wu. Deeper insights into graph convolutional networks for semi-supervised learning. *AAAI*, 2018.
- [17] Keyulu Xu, Chengtao Li, Yonglong Tian, Tomohiro Sonobe, Ken-ichi Kawarabayashi, and Stefanie Jegelka. Representation learning on graphs with jumping knowledge networks. *ICML*, 2018.

## References IV

- [18] Glen Jeh and Jennifer Widom. Scaling personalized web search. *WWW*, 2003.
- [19] Pavel Berkhin. Bookmark-coloring algorithm for personalized pagerank computing. *Internet Mathematics*, 2006.
- [20] Reid Andersen, Christian Borgs, Jennifer Chayes, John Hopcraft, Vahab S Mirrokni, and Shang-Hua Teng. Local computation of pagerank contributions. *International Workshop on Algorithms and Models for the Web-Graph*, 2007.
- [21] Oleksandr Shchur, Maximilian Mumme, Aleksandar Bojchevski, and Stephan Günnemann. Pitfalls of graph neural network evaluation. *NeurIPS Relational Representation Learning Workshop*, 2018.
- [22] Johannes Klicpera, Aleksandar Bojchevski, and Stephan Günnemann. Predict then propagate: Graph neural networks meet personalized pagerank. *ICLR*, 2019.